

This project aims at developing a machine-learning algorithm and predict if a certain mushroom is edible or poisonous by its specifications like cap shape, cap color, gill color, etc.

```
In [1]: # IMPORT ALL THE REQUIRED LIBRARY
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: #Now we will import our data
data=pd.read_csv("D:\mushrooms.csv")
```

```
In [4]: data
```

Out[4]:

| | type | cap_shape | cap_surface | cap_color | bruises | odor | gill_attachment | gill_spacing | gill_ |
|------|------|-----------|-------------|-----------|---------|------|-----------------|--------------|-------|
| 0 | p | x | s | n | t | p | f | c | |
| 1 | e | x | s | y | t | a | f | c | |
| 2 | e | b | s | w | t | l | f | c | |
| 3 | p | x | y | w | t | p | f | c | |
| 4 | e | x | s | g | f | n | f | w | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8119 | e | k | s | n | f | n | a | c | |
| 8120 | e | x | s | n | f | n | a | c | |
| 8121 | e | f | s | n | f | n | a | c | |
| 8122 | p | k | y | n | f | y | f | c | |
| 8123 | e | x | s | n | f | n | a | c | |

8124 rows × 23 columns

Data is having 8124 rows and 23 columns.

Attribute Information: (classes: edible=e, poisonous=p)

cap-shape: bell=b,conical=c,convex=x,flat=f, knobbed=k,sunken=s

cap-surface: fibrous=f,grooves=g,scaly=y,smooth=s

cap-color:
brown=n,buff=b,cinnamon=c,gray=g,green=r,pink=p,purple=u,red=e,white=w,yellow=y

bruises: bruises=t,no=f

odor: almond=a,anise=l,creosote=c,fishy=y,foul=f,musty=m,none=n,pungent=p,spicy=s

gill-attachment: attached=a,descending=d,free=f,notched=n

gill-spacing: close=c,crowded=w,distant=d

gill-size: broad=b,narrow=n

gill-color: black=k,brown=n,buff=b,chocolate=h,gray=g,
green=r,orange=o,pink=p,purple=u,red=e,white=w,yellow=y

stalk-shape: enlarging=e,tapering=t

stalk-root: bulbous=b,club=c,cup=u,equal=e,rhizomorphs=z,rooted=r,missing=? stalk-surface-
above-ring: fibrous=f,scaly=y,silky=k,smooth=s

stalk-surface-below-ring: fibrous=f,scaly=y,silky=k,smooth=s

stalk-color-above-ring:
brown=n,buff=b,cinnamon=c,gray=g,orange=o,pink=p,red=e,white=w,yellow=y

stalk-color-below-ring:
brown=n,buff=b,cinnamon=c,gray=g,orange=o,pink=p,red=e,white=w,yellow=y

veil-type: partial=p,universal=u

veil-color: brown=n,orange=o,white=w,yellow=y

ring-number: none=n,one=o,two=t

ring-type: cobwebby=c,evanescent=e,flaring=f,large=l,none=n,pendant=p,sheathing=s,zone=z

spore-print-color:
black=k,brown=n,buff=b,chocolate=h,green=r,orange=o,purple=u,white=w,yellow=y

population: abundant=a,clustered=c,numerous=n,scattered=s,several=v,solitary=y

habitat: grasses=g,leaves=l,meadows=m,paths=p,urban=u,waste=w,woods=d

```
In [5]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                  8124 non-null   object
1   cap_shape                             8124 non-null   object
2   cap_surface                           8124 non-null   object
3   cap_color                             8124 non-null   object
4   bruises                               8124 non-null   object
5   odor                                  8124 non-null   object
6   gill_attachment                       8124 non-null   object
7   gill_spacing                          8124 non-null   object
8   gill_size                             8124 non-null   object
9   gill_color                            8124 non-null   object
10  stalk_shape                           8124 non-null   object
11  stalk_root                            8124 non-null   object
12  stalk_surface_above_ring              8124 non-null   object
13  stalk_surface_below_ring              8124 non-null   object
14  stalk_color_above_ring                8124 non-null   object
15  stalk_color_below_ring                8124 non-null   object
16  veil_type                             8124 non-null   object
17  veil_color                            8124 non-null   object
18  ring_number                           8124 non-null   object
19  ring_type                             8124 non-null   object
20  spore_print_color                     8124 non-null   object
21  population                            8124 non-null   object
22  habitat                               8124 non-null   object
dtypes: object(23)
memory usage: 1.4+ MB
```

Here we can see, all the columns having object data type. As the data is object type we will use encoder to encode the value.

```
In [6]: data.isnull().sum()
```


```
Out[6]: type                0
cap_shape                  0
cap_surface                0
cap_color                  0
bruises                    0
odor                       0
gill_attachment            0
gill_spacing               0
gill_size                  0
gill_color                 0
stalk_shape                0
stalk_root                 0
stalk_surface_above_ring   0
stalk_surface_below_ring   0
stalk_color_above_ring     0
stalk_color_below_ring     0
veil_type                  0
veil_color                  0
ring_number                0
ring_type                  0
spore_print_color           0
population                 0
habitat                    0
dtype: int64
```

All the columns filled with the data, no column have null value. so we can proceed futher analysis.

We can Visualize it using heatmap.


```
In [7]: sns.heatmap(data.isnull())
```

Out[7]: <Axes: >



```
In [8]:
```

```
Out[8]: <Axes: xlabel='type', ylabel='count'>
```



```
In [9]: data.value_counts(['type'])
```

```
Out[9]: type
e      4208
p      3916
Name: count, dtype: int64
```

As our data is almost Balanced. So no need to balance the data.

```
In [10]: data.dtypes
```

```
Out[10]: type                object
cap_shape                  object
cap_surface                 object
cap_color                  object
bruises                    object
odor                      object
gill_attachment            object
gill_spacing               object
gill_size                  object
gill_color                 object
stalk_shape                object
stalk_root                 object
stalk_surface_above_ring   object
stalk_surface_below_ring   object
stalk_color_above_ring     object
stalk_color_below_ring     object
veil_type                  object
veil_color                 object
ring_number                object
ring_type                  object
spore_print_color          object
population                 object
habitat                    object
dtype: object
```

here we can see our data type is objective we need to encode the data before going for futher steps of analysis.

```
In [11]: from sklearn.preprocessing import LabelEncoder
def label_encoded(x):
    le = LabelEncoder()
    le.fit(x)
    print(x.name, le.classes_)
    return le.transform(x)
```

Here we used made the funtion for encoding out data. We used LabelEncoder for this.

```
In [12]: for col in data.columns:
        data[str(col)] = label_encoded(data[str(col)])
```

```
type ['e' 'p']
cap_shape ['b' 'c' 'f' 'k' 's' 'x']
cap_surface ['f' 'g' 's' 'y']
cap_color ['b' 'c' 'e' 'g' 'n' 'p' 'r' 'u' 'w' 'y']
bruises ['f' 't']
odor ['a' 'c' 'f' 'l' 'm' 'n' 'p' 's' 'y']
gill_attachment ['a' 'f']
gill_spacing ['c' 'w']
gill_size ['b' 'n']
gill_color ['b' 'e' 'g' 'h' 'k' 'n' 'o' 'p' 'r' 'u' 'w' 'y']
stalk_shape ['e' 't']
stalk_root ['?' 'b' 'c' 'e' 'r']
stalk_surface_above_ring ['f' 'k' 's' 'y']
stalk_surface_below_ring ['f' 'k' 's' 'y']
stalk_color_above_ring ['b' 'c' 'e' 'g' 'n' 'o' 'p' 'w' 'y']
stalk_color_below_ring ['b' 'c' 'e' 'g' 'n' 'o' 'p' 'w' 'y']
veil_type ['p']
veil_color ['n' 'o' 'w' 'y']
ring_number ['n' 'o' 't']
ring_type ['e' 'f' 'l' 'n' 'p']
spore_print_color ['b' 'h' 'k' 'n' 'o' 'r' 'u' 'w' 'y']
population ['a' 'c' 'n' 's' 'v' 'y']
habitat ['d' 'g' 'l' 'm' 'p' 'u' 'w']
```

```
In [13]: data.head()
```

Out[13]:

| | type | cap_shape | cap_surface | cap_color | bruises | odor | gill_attachment | gill_spacing | gill_size |
|---|------|-----------|-------------|-----------|---------|------|-----------------|--------------|-----------|
| 0 | 1 | 5 | 2 | 4 | 1 | 6 | 1 | 0 | |
| 1 | 0 | 5 | 2 | 9 | 1 | 0 | 1 | 0 | |
| 2 | 0 | 0 | 2 | 8 | 1 | 3 | 1 | 0 | |
| 3 | 1 | 5 | 3 | 8 | 1 | 6 | 1 | 0 | |
| 4 | 0 | 5 | 2 | 3 | 0 | 5 | 1 | 1 | |

5 rows × 23 columns

```
In [14]: data.columns
```

Out[14]: Index(['type', 'cap_shape', 'cap_surface', 'cap_color', 'bruises', 'odor', 'gill_attachment', 'gill_spacing', 'gill_size', 'gill_color', 'stalk_shape', 'stalk_root', 'stalk_surface_above_ring', 'stalk_surface_below_ring', 'stalk_color_above_ring', 'stalk_color_below_ring', 'veil_type', 'veil_color', 'ring_number', 'ring_type', 'spore_print_color', 'population', 'habitat'], dtype='object')


```
In [15]: data.corr()
```

Out[15]:

| | type | cap_shape | cap_surface | cap_color | bruises | odor |
|-----------------|-----------|-----------|-------------|-----------|-----------|-----------|
| type | 1.000000 | 0.052951 | 0.178446 | -0.031384 | -0.501530 | -0.093552 |
| cap_shape | 0.052951 | 1.000000 | -0.050454 | -0.048203 | -0.035374 | -0.021935 |
| cap_surface | 0.178446 | -0.050454 | 1.000000 | -0.019402 | 0.070228 | 0.045233 |
| cap_color | -0.031384 | -0.048203 | -0.019402 | 1.000000 | -0.000764 | -0.387121 |
| bruises | -0.501530 | -0.035374 | 0.070228 | -0.000764 | 1.000000 | -0.061825 |
| odor | -0.093552 | -0.021935 | 0.045233 | -0.387121 | -0.061825 | 1.000000 |
| gill_attachment | 0.129200 | 0.078865 | -0.034180 | 0.041436 | 0.137359 | -0.059590 |
| gill_spacing | -0.348387 | 0.013196 | -0.282306 | 0.144259 | -0.299473 | 0.063936 |
| gill_size | 0.540024 | 0.054050 | 0.208100 | -0.169464 | -0.369596 | 0.310495 |
| gill_color | -0.530566 | -0.006039 | -0.161017 | 0.084659 | 0.527120 | -0.129213 |
| stalk_shape | -0.102019 | 0.063794 | -0.014123 | -0.456496 | 0.099364 | 0.459766 |

```
In [ ]: We can Visualize the correlation using Heatmap. as follow
```

```
In [16]: plt.figure(figsize=(20,15))
sns.heatmap(data.corr(), annot=True, linewidths=1, linecolor='black',fmt=' .2f'
```



Here we can see vell-type have no effect on our target column, so we can drop that column.

```
In [17]: data.head(10)
```

Out[17]:

| | type | cap_shape | cap_surface | cap_color | bruises | odor | gill_attachment | gill_spacing | gill_size |
|---|------|-----------|-------------|-----------|---------|------|-----------------|--------------|-----------|
| 0 | 1 | 5 | 2 | 4 | 1 | 6 | 1 | 0 | |
| 1 | 0 | 5 | 2 | 9 | 1 | 0 | 1 | 0 | |
| 2 | 0 | 0 | 2 | 8 | 1 | 3 | 1 | 0 | |
| 3 | 1 | 5 | 3 | 8 | 1 | 6 | 1 | 0 | |
| 4 | 0 | 5 | 2 | 3 | 0 | 5 | 1 | 1 | |
| 5 | 0 | 5 | 3 | 9 | 1 | 0 | 1 | 0 | |
| 6 | 0 | 0 | 2 | 8 | 1 | 0 | 1 | 0 | |
| 7 | 0 | 0 | 3 | 8 | 1 | 3 | 1 | 0 | |
| 8 | 1 | 5 | 3 | 8 | 1 | 6 | 1 | 0 | |
| 9 | 0 | 0 | 2 | 9 | 1 | 0 | 1 | 0 | |

10 rows × 23 columns

```
In [18]: data=data.drop(['veil_type'],axis=1)
```

```
In [19]: data.head(5)
```

Out[19]:

| | type | cap_shape | cap_surface | cap_color | bruises | odor | gill_attachment | gill_spacing | gill_size |
|---|------|-----------|-------------|-----------|---------|------|-----------------|--------------|-----------|
| 0 | 1 | 5 | 2 | 4 | 1 | 6 | 1 | 0 | |
| 1 | 0 | 5 | 2 | 9 | 1 | 0 | 1 | 0 | |
| 2 | 0 | 0 | 2 | 8 | 1 | 3 | 1 | 0 | |
| 3 | 1 | 5 | 3 | 8 | 1 | 6 | 1 | 0 | |
| 4 | 0 | 5 | 2 | 3 | 0 | 5 | 1 | 1 | |

5 rows × 22 columns

Since the outcome has only two variable we will use binary classification model.

```
In [20]: # Now split data in dependent and independent variables
x=data.iloc[:,1:]
y=data.iloc[:,-22]
```

```
In [21]: x.shape
```

Out[21]: (8124, 21)

```
In [22]: y.shape
```

Out[22]: (8124,)

```
In [23]: y.head(10)
```

Out[23]:

| | |
|---|---|
| 0 | 1 |
| 1 | 0 |
| 2 | 0 |
| 3 | 1 |
| 4 | 0 |
| 5 | 0 |
| 6 | 0 |
| 7 | 0 |
| 8 | 1 |
| 9 | 0 |

Name: type, dtype: int32

```
In [24]: from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
lr=LogisticRegression()
from sklearn.metrics import accuracy_score, confusion_matrix, classification_r
```

Here we have imported all required training models, we check the best fit model.

```
In [25]: for i in range(0,1000):
    x_train, x_test, y_train, y_test= train_test_split(x,y,random_state= i,test_s
    lr.fit(x_train,y_train)
    pred_train= lr.predict(x_train)
    pred_test= lr.predict(x_test)
    if round (accuracy_score(y_train,pred_train)*100,1)==round(accuracy_score(
        print("At Random state",i, "The model perform very well")
        print("At random State:",i)
        print("Training r2_score",accuracy_score(y_train,pred_train)*100)
        print("testing r2 score ",accuracy_score(y_test,pred_test)*100)
```

At Random state 10 The model perform very well
At random State: 10
Training r2_score 95.24542237267272
testing r2 score 95.19999999999999
At Random state 18 The model perform very well
At random State: 18
Training r2_score 95.30697030312356
testing r2 score 95.26153846153846
At Random state 28 The model perform very well
At random State: 28
Training r2_score 94.99923065086936
testing r2 score 95.01538461538462
At Random state 34 The model perform very well
At random State: 34
Training r2_score 95.33774426834898
testing r2 score 95.26153846153846
At Random state 36 The model perform very well
At random State: 36
Training r2_score 94.8915217725804
testing r2 score 94.80000000000000
At random State: 210

Training r2_score 95.56854900753962

testing r2 score 95.56923076923077

At random state 210, we are get better r2 score.

```
In [26]: x_train, x_test, y_train, y_test= train_test_split(x,y,random_state=210,test_s
```

```
In [27]: from sklearn.metrics import classification_report
print(classification_report(y_test,pred_test))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.52 | 0.53 | 0.53 | 844 |
| 1 | 0.48 | 0.46 | 0.47 | 781 |
| accuracy | | | 0.50 | 1625 |
| macro avg | 0.50 | 0.50 | 0.50 | 1625 |
| weighted avg | 0.50 | 0.50 | 0.50 | 1625 |


```
In [28]: pred_lr = lr.predict(x_test)
from sklearn.model_selection import cross_val_score
lss=accuracy_score(y_test,pred_lr)
for j in range(2,10):
    lsscore= cross_val_score(lr,x,y,cv=j)
    lsc=lsscore.mean()
    print("at cv:-", j)
    print("Cross Validation scre is:-",lsc*100)
    print("Accuracy Score:-", lss*100)
    print("\n")
```

```
at cv:- 2
Cross Validation scre is:- 81.29000492368291
Accuracy Score:- 94.95384615384616
```

```
at cv:- 3
Cross Validation scre is:- 81.25307730182176
Accuracy Score:- 94.95384615384616
```

```
at cv:- 4
Cross Validation scre is:- 82.87789266371246
Accuracy Score:- 94.95384615384616
```

```
at cv:- 5
Cross Validation scre is:- 83.93488442591892
Accuracy Score:- 94.95384615384616
```

```
at cv:- 6
Cross Validation scre is:- 87.87543082225504
Accuracy Score:- 94.95384615384616
```

```
at cv:- 7
Cross Validation scre is:- 88.36851830636915
Accuracy Score:- 94.95384615384616
```

```
at cv:- 8
Cross Validation scre is:- 86.70541047670766
Accuracy Score:- 94.95384615384616
```

```
at cv:- 9
Cross Validation scre is:- 87.418672586001
Accuracy Score:- 94.95384615384616
```

```
at cv:- 3

Cross Validation scre is:- 81.12998522895126

Accuracy Score:- 95.19999999999999
```

```
In [29]: lssscore_selected=cross_val_score(lr,x,y,cv=3).mean()
print("The cv score is: ", lssscore_selected,"\nThe accuracy score is: ", lss)
```

```
The cv score is: 0.8125307730182176
The accuracy score is: 0.9495384615384616
```

```
In [30]: from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
```

```
In [31]: Model= [LogisticRegression(),RandomForestClassifier(),GaussianNB(),DecisionTre

for m in Model:
    m.fit(x_train,y_train)
    m.score(x_train,y_train)
    predm=m.predict(x_test)
    print('Accuracy score of', m, 'is:')
    print (accuracy_score(y_test,predm))
    print (confusion_matrix(y_test,predm))
    print (classification_report(y_test,predm))
    print('\n')
```

Accuracy score of LogisticRegression() is:
0.9556923076923077
[[804 40]
[32 749]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.95 | 0.96 | 844 |
| 1 | 0.95 | 0.96 | 0.95 | 781 |
| accuracy | | | 0.96 | 1625 |
| macro avg | 0.96 | 0.96 | 0.96 | 1625 |
| weighted avg | 0.96 | 0.96 | 0.96 | 1625 |

Accuracy score of RandomForestClassifier() is:
1.0
[[844 0]
[0 781]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 844 |
| 1 | 1.00 | 1.00 | 1.00 | 781 |
| accuracy | | | 1.00 | 1625 |
| macro avg | 1.00 | 1.00 | 1.00 | 1625 |
| weighted avg | 1.00 | 1.00 | 1.00 | 1625 |

Accuracy score of GaussianNB() is:
0.9243076923076923
[[774 70]
[53 728]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.92 | 0.93 | 844 |
| 1 | 0.91 | 0.93 | 0.92 | 781 |
| accuracy | | | 0.92 | 1625 |
| macro avg | 0.92 | 0.92 | 0.92 | 1625 |
| weighted avg | 0.92 | 0.92 | 0.92 | 1625 |

Accuracy score of DecisionTreeClassifier() is:
1.0
[[844 0]
[0 781]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 844 |
| 1 | 1.00 | 1.00 | 1.00 | 781 |
| accuracy | | | 1.00 | 1625 |
| macro avg | 1.00 | 1.00 | 1.00 | 1625 |
| weighted avg | 1.00 | 1.00 | 1.00 | 1625 |

Accuracy score of SVC() is:
0.9932307692307693
[[844 0]
[11 770]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 0.99 | 844 |
| 1 | 1.00 | 0.99 | 0.99 | 781 |
| accuracy | | | 0.99 | 1625 |
| macro avg | 0.99 | 0.99 | 0.99 | 1625 |
| weighted avg | 0.99 | 0.99 | 0.99 | 1625 |

accuracy score [0.9556923076923077, 1.0, 0.9243076923076923, 1.0, 0.9932307692307693]
RandomForestClassifier()=1.0

DecisionTreeClassifier() =1.0

Here all the three performing well so we can save any one model.


```
In [37]: dtc= DecisionTreeClassifier()  
dtc.fit(x_train,y_train)  
dtc.score(x_train,y_train)  
  
dtcpred=dtc.predict(x_test)  
print(accuracy_score(y_test,dtcpred ))  
print(confusion_matrix(y_test,dtcpred ))  
print(classification_report(y_test,dtcpred ))
```

```
1.0  
[[844  0]  
 [ 0 781]]  
  
              precision    recall  f1-score   support  
  
    0               1.00      1.00      1.00     844  
    1               1.00      1.00      1.00     781  
  
 accuracy               1.00     1625  
 macro avg              1.00     1625  
weighted avg              1.00     1625
```

AUC ROC Curve:

```
In [38]: from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds= roc_curve(dtcpred, y_test)
roc_auc= auc(fpr, tpr)

plt.figure()
plt.plot(fpr,tpr, color='darkorange', lw=10,label='ROC curve (area= %0.2f)' %r
plt.plot([0,1],[0,1],color ='navy', lw=10, linestyle= '--')
plt.xlim([0.0, 1.0])
plt.xlim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
```



Now Saving best Model

```
In [39]: import pickle
filename='Mushroom_Pred.pkl'
pickle.dump(dtc,open(filename,'wb'))
```

Conclusion

```
In [40]: import numpy as np
a=np.array(y_test)
predicted= np.array(dtc.predict(x_test))
df_com= pd.DataFrame({'original':a, 'predicted':predicted}, index= range(len(a)
df_com
```

Out[40]:

| | original | predicted |
|------|----------|-----------|
| 0 | 1 | 1 |
| 1 | 0 | 0 |
| 2 | 0 | 0 |
| 3 | 0 | 0 |
| 4 | 0 | 0 |
| ... | ... | ... |
| 1620 | 0 | 0 |
| 1621 | 0 | 0 |
| 1622 | 0 | 0 |
| 1623 | 0 | 0 |
| 1624 | 0 | 0 |

1625 rows × 2 columns

```
In [ ]:
```